Assessing quality of music generated by the Music Transformer with an American Folk dataset

github.com/gregwinther/folk_transformer and https://gregwinther.github.io/folk_transformer/

Sebastian G. Winther-Larsen

Center for Computing in Science Education, Department of Physics, University of Oslo

Tom F.Hansen

Institute of Informatics, University of Oslo

Bjørn Iversen

Institute of Informatics, University of Oslo

Abstract—Since the publication of the music transformer in 2017, music generators based on machine learning (ML) with the transformer architecture has been state of the art, making it possible to generate realistic music for over a minute. Up to now the results have mainly been succesful for classical piano music, but have failed for more technically complex music such as jazz. Americana folk music can be seen as a stepping stone to more complex music. In this work we generate Americana folk music with 2 different approaches; transfer learning from the MAESTRO dataset and solely training on a dataset of Americana music. Moreover, we propose a framework for the evaluation of these results in a quantitative and qualitative way, leaning on Evidence-Based Design for Assessment and Evaluation. Results from the quality assessment indicate that both models have been moderately succesful, and the model based on transfer learning is favoured. In a rating process people have rated the songs to in average 2.6, where 1 is not realistic and 5 is realistic.

Index Terms—music generation, transformer model, evaluation, transfer learning, Americana

I. MOTIVATION AND INTRODUCTION

The attention-based transformer model [1], is today recognised as the best performing sequential machine learning model, surpassing RNN-based models in most cases, mainly because of its superior memory efficiency and shorter training because it is parallelizable. While originally applied to natural language processing (NLP), which today has mature implementations, the architecture can also be applied for other sequential learning tasks, such as music generation [2]. The music transformer developed in the Magenta project is trained on the MAESTRO dataset [3]. By setting a primer – a start music sequence, the model generates new music with good results along the same lines as the training set. With other primers than the regular and systematic classical music in MAESTRO, the quality of the output is varying.

A **main goal** of this work is to describe a framework for the quality assessment of different approaches to the generation of music with the transformer model. This is further motivated by the attempt to generate, and hence evaluate, more varied music than up to now has been successfully generated. We compare two different Transformer implementations and use our proposed evaluation framework to assess the quality of each approach. At the core of such an evaluation framework is an evaluation format that combines a form of quantitative and qualitative evaluation technique. In ML-modelling dealing with a creative art like music, the subjective experience of the result will be an important part in the evaluation, hence the importance of increasing the evaluation scheme from the purely technical approach that is normal for most ML-models.

To this end, we wish to employ the transformer music model to a subgenre of music to which such a model has not been extensively applied. Initial findings from applying the transformer to jazz music has shown some limitations [4], applying LSTM networks to Blues has been moderately succesful [5] and applying the transformer model to pop piano music seems to work well [6]. This may be due to the rules and structure of pop and classical music - classical music often has formal rules, the epitome of which is the fugue [7]; and pop music follows some very clear norms [8]. While even Free Jazz has *some* rules, it readily falls into the category of the type of rhythmic music with the least amount of structure, that it is "characterized by the absence of set chord patterns or time patterns"[9].

American roots music, encompassing spirituals, cajun music, cowboy music, work songs, but also early blues such as Dixieland; from now on referred to as "Americana", presents itself as hitherto unexplored territory. It also provides a nice stepping stone towards more "unstructured" music as it often allows for improvisation, but otherwise retains a relatively rigid structure [10]. We have therefore collected a dataset of MIDI files of Americana music, which we will use in our quality assessment framework. A **secondary goal** is therefore to realistically generate Americana music with the transformer architecture. A natural **third goal** in the process-steps is to investigate the superiority of music generators based on a significantly higher number of songs, even though this could be music from other genres, by utilizing transfer learning from the MAESTRO dataset in the music transformer [2].

In the next chapter (II) we present related work that we have elaborated on, both dealing with music generation and music evaluation. Chapter III deals with methods used for evaluation and modelling. In chapter IV we describe the 2 musical datasets. Chapter V and VI describes results and discussion, before chapter 7 illustrate challenges and sum up conclusions.

II. RELATED WORK

The transformer model is considered state of the art in music generation, surpassing RNN-based models in the last few years. Both are sequential models, but the attention principle at the core of the transformer facilitates remembering coherence over longer sections of sequences and highlights especially important sections. From generating music of 10s of seconds with RNNs, it is now possible to generate a minute of coherent realistic music [2]. Still, there are a lot of unresolved challenges, like generating long sections over some minutes, in highly irregular compositions and multi-channel or many-instrumental signals. To combat these challenges the improvement of the transformer model has a high focus in the research community. Some of the most recent attempts are the Transformer-XL [11] and the Reformer [12]. The Reformer claims to be trained on a standard computer with a single GPU. In this analysis, we will utilize and adaption of the Transformer-XL architecture for music.

A. Music transformers

In the existing papers considering music generation Transformer models we want to emphasize some important related works besides the aforementioned music transformer, built on the classical piano music dataset MAESTRO. These works are somewhat diversified in different music genres, describing the span of existing music genres generated by the transformer, "setting the stage" for the Americana music in this work. In the pop music transformer [6] pop piano music is generated by the transformer, a setting quite similar to the classical piano music in the music transformer [2].

Increasing the complexity of music generated, Wu & Yang describes the Jazz transformer which is in the other end of the spectrum related to complexity. Here the Transformer-XL architecture is utilized to model lead sheets of jazz music. Moreover, the model endeavours to incorporate structural events present in the Weimar Jazz Database (WJazzD) for inducing structures in the generated music, the results are not impressive according to subjective listening tests and technical evaluation based on pitch class histogram entropy. Listening tests shows a clear gap between the ratings of the generated and real compositions. The work analyses the missing parts and presents a prediction system which analytically shed light on why machine-generated music to date still falls short of the artwork of humans. This includes analyzing the statistics of the pitch class, grooving, and chord progression, assessing the structures of the music with the help of the fitness scape plot, and evaluating the model's understanding of Jazz music. The evaluation scheme and the failure to generate such complex music is relevant in the generation of Americana.

In a system called LakhNES, multi-instrumental music is generated with a Transformer [13]. Their success of music generation with the piano score generation is partially explained by the large volumes of symbolic data readily available for that domain. They leverage the recently-introduced NES-MDB dataset of four-instrument scores from an early video game sound synthesis chip[1]. They found this data to be

[1]The Nintendo Entertainment System (NES)

well-suited to training with the Transformer architecture. The model was further improved with a pre-training technique to leverage the information in a large collection of heterogeneous music, namely the Lakh MIDI dataset. By performing transfer learning on the NES-MDB dataset, both the qualitative and quantitative performance from the target dataset was significantly improved. The rare use of transfer learning in music generation with the transformer is a relevant foundation for use of transfer learning in this work.

Gan *et al.* use the transformer architecture to generate music, but with another approach. In a system called Foley Music, they synthesize music from a silent video about people playing instruments [14]. A relationship between body keypoints and classical MIDI recordings is established. Music generation is then formulated as a motion-to-MIDI translation problem, represented with a graph transformer framework that predicts MIDI from motion. By testing the generator on different music performances the results are proven to outperform several existing systems in music generation, again for classical music, but now not only piano music.

Summarized there are a few examples of successfully generating music with the transformer, mainly with classical music. However, there is little work on generating intentionally the same music generator with different approaches, as will be the case in our work. Another new approach is to use transfer learning from an existing high-performance music model, opposed from the "simpler" NES-MDB dataset.

B. Evaluation of AI-generated music

In the evaluation of music generation models based on machine learning we would like to point out the works by Yang & Lerch [15] describing the technical evaluation system mgeval, utilized in this work, Wu & Yang [4] describing a system of qualitative listening tests and 5 objecive measurements of the structure in the music, and lastly the qualitative evaluation measures carried out by Sturm & Ben-Tal [16], focusing both on qualitative interviews and statistics about the generated notes. These works describe qualitative and quantiative evaluations between **real music and a generated music**. In this work we elaborate on these works and combine approaches in a proposal of a common framework for evaluating **different** transformer models for music generation.

III. METHODS

Acting as a base and for exemplification of the benchmark architecture, Americana music is generated in 2 different model concepts:

1) Utilize transfer learning with MAESTRO music transformer as a base, and train with the full dataset of Americana music.
2) Train a new transformer model only with the full Americana dataset.

We hypothesise that the transfer learning model will result in the best performing model, but an important issue is what makes up the best model and how to evaluate such a subjective "sequence-result" as music in a fair and trustworthy manner. The objective is to investigate if the transfer learning model

is performing significantly better than a model only trained on a single dataset, such as has been shown in other ML applications, like image classification [17], even though the MAESTRO dataset is totally different from the Americana dataset.

An attempt to sort this out is by evaluating in a quantitative and qualitative way, as will be the charecteristics in our work. The quantitative, hence objective, way can shortly be described as a technical comparison of the predicted signal and the real signal. Principles by Yang & Lerch [15] and Wu & Yang [4] will be utilized, further described in the section below.

The qualitative part constitutes a music expert judgement, based on listening to the generated music files from the objective evaluation. In a second, and survey-based part, 12 random persons are asked to rate the different music files.

A. Qantitative evaluation

For quantitative evaluation, we will be using the objective evaluation toolbox `mgeval` [15]. The toolbox extracts absolute metrics from MIDI files, which allows for inspection of properties of the dataset used for training and the generated dataset. The features extracted for absolute measures are divided into pitch-based features:

- Pitch count (PC)
- Pitch class histogram (PCH)
- Pitch class transition matrix (PCTM)
- Pitch range (PR)
- Average pitch interval (PI)

and rhythm-based features:

- Note count (NC)
- Average inter-onset-interval (IOI)
- Note length histogram (NLH)
- Note length transition matrix (NLTM)

These metrics can then be used to acquire the relative metrics between datasets with the use of exhaustive cross-validation to acquire the distance between each sample of the same set (intra-dataset) and another set (inter-dataset).

B. Evidence-Based Design for Assessment and Evaluation

As a rigorous and well-proven approach do construct a framework for assessing music composed by artificial intelligence models, we propose adapting the methodology of Evidence-Centered Design [18, 19]. By working with Evidence-Centered Design (ECD), we engage in a modern approach to assessment design, and for assessing complex knowledge and practices. ECD is originally applied to the contruction of psychometric learning assessment tools. Through to completion, it would take several years to construct such a tool, something that is well outside the scope of this study. However, we prospose to begin with the first step within ECD - **Domain Analysis**. This involves exploratory interviews of experts in the field in order to contruct a thematically organized and prioritized list of knowledge and practices to assess. Specific to our study, we find it necessary to talk to professional musicians and composers in order to uncover what actually makes a good composition. The method described has similarities with studies performed by Sturm & Ben-Tal [16].

To supplement the qualitative evaluation a survey was distributed to random people with unknown backgrounds. The participants were presented to 5 generated songs with length between 40 seconds and one minute, facilitated in a Google survey form. The songs were generated by sending a primer of the first three bars (five seconds) of five different original songs from the Americana dataset to the trained model. Each song were presented in two versions, one generated by Transfer-Americana and one by Americana. The participants had to choose the one they liked best in each pair of songs. In addition they had to rate the best version related to how realistic it was compared to human generated music. The range was 1 (worst) to 5 (best).

C. Modelling process

MIDI-files was first preprocessed to fit into the format used by the transformer model. This was carried out by utilizing an implementation of a transformer-XL [11], adapted for music, incorporated in a package called Musicautobot [20]. Scripts for transfer learning and single learning were built around this package. Both models had the same model topology, used the same tuning parameters, and were trained on four NVIDIA 2080TI GPUs at University of Oslo for approximately ten hours each. The models had the following parameters, `n_layers`=16 `d_model` = 512, `d_head` = 64, `n_head` = 8, `d_inner` = 2048, `mem_len` = 512, and `batchsize` = 8.

IV. DATASETS

A brief summary of each of the datasets we have used in this study can be found in Table I.

MAESTRO [3] (MIDI and Audio Edited for Synchronous TRacks and Organization) is a dataset with over 200 hours of virtuosic piano perfomances captured with a fine alignment of approximately 3ms between note labels and audio waveforms.

The data is a product of performances in the International Piano-e-competition. During each installment of the competiton, virtuoso pianists perform on Yamaha Disklaviers which, in addition to being concert-quality acoustic grand pianos, utilize integrated high-precision MIDI capture and playback.

Since the MAESTRO dataset contains MIDI recordings from competions, the pieces are from a select set for each year. This means that many of the pieces are the same, but may includes much variation within each performer's interpretation of the piece.

The Americana dataset is constructed from musical scores by Benjamin Robert Tubb [2]. These scores are in the public domain and composed between the ealry 1800s and 1922. The genres range from blues, ragtime, naval songs, hyms, minstrel songs and spirituals.

V. RESULTS

A. Quantitative evaluation

For this evaluation we have picked ten melodies from both the Americana dataset and the MAESTRO dataset, as well as using ten different primers to generate ten melodies with

---

[2]sourced online

Table I
DATA SET DESCRIPTION.

	MAESTRO v2	Americana
No. of songs	1282	5711
Total time [hours]	201	329
Mean length [min]	9.4	3.45

Table II
RESULTS FROM RELATIVE MEASUREMENTS OF INTRA-SET DISTANCES.

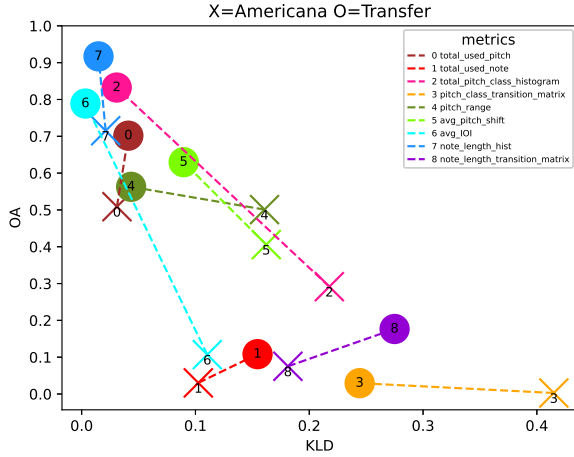|	Training		Americana		Transfer	
	Intra-set		Intra-set		Intra-set	
	mean	STD	mean	STD	mean	STD
PC	11.777	1.700	3.577	0.636	6.466	0.706
NC	442.377	73.911	11.377	1.542	21.955	4.678
PCH	0.384	0.013	0.164	0.007	0.365	0.026
PCTM	821.396	56.336	10.475	0.292	25.991	3.303
PR	19.088	1.638	5.266	0.611	10.400	2.038
PI	3.836	0.510	0.848	0.128	2.032	0.479
IOI	0.046	0.006	0.105	0.007	0.235	0.022
NLH	0.471	0.027	0.263	0.022	0.466	0.043
NLTM	316.896	23.442	23.241	1.359	46.778	4.027

Figure 1. Visualisation of the KLD and OA of the Intra-set generated from Training set and the American/Transfer sets.


Figure 2. Absolute measurements of average pitch class transition matrix (PCTM) and note length transition matrix (NLTM) from the Americana training set, Americana generated, transfer generated and MAESTRO training set (sect IV.A.)
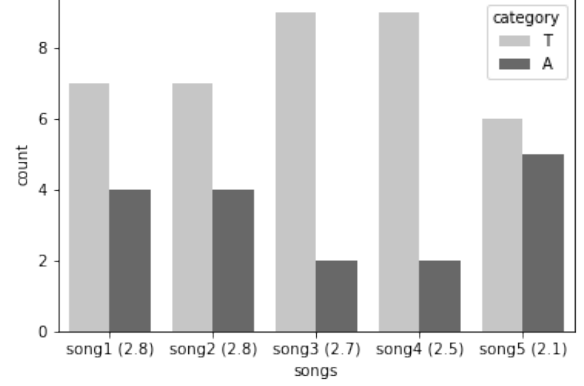

Figure 3. Bars show a count of the number of votes, respectively for Transfer-Americana and Americana, for each of the 5 songs. The numbers in parenthesis below the Bars is the average rating for each of the songs, when compared to a realistic human made song (1 is worst, 5 is best)

the Americana trained and the transfer trained melody sets. The first five melodies, denoted 0-4, is generated by the same primers as in the survey melodies visualized in Figure 3.

The results for the absolute measures of note length transition matrix in Figure 2 shows us that the Americana/training dataset have a greater variety in note length than the MAE-STRO set (comparing similarity between column 1 and 4). We can see that the melodies generated with the Americana dataset have similar variety and that the melodies generated with the Transfer set influenced by the MAESTRO set preservers most variety (comparing column 2 and 3 with the real melodies in column 1). The result for pitch class transition matrices on the other hand shows a decrease in variety in the transfer generated melodies.

The results for the relative measurements shown in Table II and Figure 1. From the comparisons we can see that both the Americana and Transfer sets have notably smaller mean values in the Intra-set across the board, apart from average inter-onset-interval (IOI). This imply that there is less variation in the generated samples than the training set, while the time between each note also increases as seen in the IOI. We can also see that the transfer set have twice as high mean than the Americana set which indicate notable improvements in the generated melodies, however the standard deviation increases as well which imply that the mean in the transfer set is less reliable.

Lastly, we will look at the Kullback–Leibler divergences (KLD) and Overlapping area (OA) for the Inter-sets generated for training-Americana and training-transfer as seen in Figure 1. We can see that the OA for the transfer set is greater for all the metrics which again indicates improvements in transfer over the Americana set. However, there is also an increase in the KLD for pitch count, note count and note length, which indicates that these metrics are less reliable.

B. Qualitative evaluation

1) Survey results: Figure 3 shows the count of the number of votes, respectively for Transfer-Americana and Americana,

Figure 4. Boxplot of ratings for each of the 5 songs when compared to a realistic human made song (1 is worst, 5 is best). Median in horisontal black line, box defines by 25-75 percentile. The red dots are all registered scores.

for each of the 5 songs. The numbers in parenthesis below the Bars is the average rating for each of the songs when compared to a real human-made song (1 is worst, 5 is best). Inspecting the figure we see a general trend where Transfer-Americana is favoured in 4 of 5 songs. Song 5 has a closer score and is also the song with the lowest average rating of 2.1. 69 % of all choices favour Transfer-Americana. A paired t-test was conducted for this result, describing a significant favouring of Transfer-Americana, with a p-value of 0.07. These results verified the hypothesis that a music transformer model will perform better when it is based on the pre-trained model, increasing the total number of music samples, even though the base model was trained on the classical MAESTRO dataset. It must be pointed out the uncertainty related to the relatively low number of samples in the survey. In Figure 4 a boxplot is visualizing the single scorings in a more detailed manner. Red points mark all the single scorings. The median for song 1 through 4 has a score of 3 and also single scorings of 4 and 5, indicating that some people consider the songs as comparable to human-made music.

In this survey some people have made comments about the overall performance, focusing more on the subjective experience. A common comment were that the songs appear to all get worse the longer they go on, as in the notes get more sparse and less cohesive. The songs also seem to go all the way through a scale. Such patterns make it very clear it is not a trained human playing Another frequent comment was that it might have been better if the songs from the different generators were shuffled for each question so the listeners didn't naturally consider the differences between the two. This issue could color each opinion. Lastly it was the meaning of each rate in the scoring system was unclear related the human level of music composition.

2) Expert Interviews: Several samples of generated music were played for three professional musicians and composers. All of these have degrees in either music performance or composition, as well as experience with songwriting as well as performing music within a few of the subgenres of Americana, namely blues and folk. None of the musicians have any exprience or great knowledge of methods from ML or AI.

After a brief introduction to the concept of AI-generated music, all of the musicians volunteered genres that would be easiest for a computer to generate. Classical music, especially baroque, blues and pop were all mentioned, as music within these genres "have specific rules" and that "there exists a recipe" for these genres. When asked if they could name genres that they thought would be difficult to model, jazz was the dominant response. This initial hunch is interesting because it coincides with the space of genres that span the current work on music in AI.

The concensus amongst the musicians were that the music generated from the transfer-based model was best compared with the other model, but they all emphasised that in absolute terms the generated music was not to be considered good.

From the comments, some general guidelines for music assessment can be extracted. Broadly summed up, one would like a musical piece to "follow the rules" both when it comes to rythm and scale, but also to "break the rules" enough to make the piece interesting, but not chaotic. Preservation of long-term structure is important. "It used the same stuff again!", was loudly exclaimed by one of the musicians after a musical theme was repeated in one of the pieces.

Regarding scales and pitch variation, some interesting comments were made that exemplifies a common theme amongst our panel of experts; "It sounds like a child, who has just learned a new scale, trying to remember which notes are correct by testing different ones." The musician that made this comment was asked to elaborate and said that it is okay to venture outside the designated scale of the composition to create suspense, but that the generated music tended to do this very suddenly and without any transtion to the new regime. Moreover, "a structural break can be nice to capture the attention of the listener, but don't overdo it". The musicians quickly adapted to a certain listening mindset, and started counting "errors" the models made; "in this piece there was only three scale breaches. Not too bad." Some musicians pointed to sections in the music that appeared to them as "correct" structural breaks, but though that the model seemed "confues itself" by introducing them.

The musician thought that the models performance regarding rythm was wide-ranging, but seemed better than the "scale rule-breaking". Some examples of grave rythmic errors were pointed out, for example "that dotted note was incredibly weird!" and "what a very unexpected long rest".

The results from the interview session is much in line with results described by Sturm & Ben-Tal [16], collecting this kind of information from a website presenting generated music with RNN-based models. However, an important difference is the lack of desription of music background from the evaluators in that study.

VI. DISCUSSION

Our goal nr 1 was to propose a framework for the evaluation of music generated by machine learning with a Transformer architecture. The combination of a quantitative system and two approaches to qualitative evaluation is undoubtedly a step

in this direction, giving a broader view of the results, also touching upon the more subjective parts of an evaluation. We find that the guidelines provided by the musicians and the survey were in line with what the quantitative evaluation toolbox `mgeval` [15] is trying to do. However, we emphasise the need for such quantitative tools to have the basis in an expert opinion as suggested by the ECD framework [18, 19].

A naturally related goal nr 2 of the work was to verify that a model trained by transfer learning and therefore based on more samples will perform better, even though the base model constitutes of songs of a different music genre, and as goal nr 3 show that it is possible to generate more complex music than generated in earlier related work. Results from qualitative and quantitative evaluation clearly show that Transfer-Americana perform better than Americana, fulfilling goal 2. Scorings from comparison with human-made music to some degree supports goal 3, with an average rating of 2.6 of 5, but there is still a way to go. Especially for the songs from the Americana model, we see that the melodic part degrades faster than for Transfer-Americana, a pattern which probably has affected the comparisons and the scoring.

An interesting pattern is the variation in count of superiority between songs in figure 3. In song 1-4, the count shows a clear overweight on Transfer-Americana, especially evident for song 3 and 4. For song 5 the count is close to similar. When we compare this to the technical evaluation there is support in this pattern for song 1-4. For song 3 and 4 there is a longer distance between Americana and Transfer-Americana and the difference between the 2 versions should be higher. It could be that the composition and complexity in certain primers are more dependent on a bigger database of trained songs, than other primers. This is something to further analyse in future work. Scorings in figure 4 return many 3's and some 4's and 5's, but also 1's, indicating that some people think this is truly realistic music, while others condsider human-made and generated music as two different worlds.

Processing of data and adapting a framework to fit our problem took an extraordinary amount of time. Optimally the models should have been trained for a longer period, and the resulting model is a minimal viable model. Transformer models can be monstrous an unwieldy structures, illustrated by this quote from Vaswani *et al.*: "The Transformer allows for significantly more parallelization and can reach a new state of the art in translation quality after being trained for as little as twelve hours on eight P100 GPUs". We find it odd to assume that it is normal to have access to a rack of eight professional-grade GPUs [1], illustrating the importance of further development of the transformer to increase the efficiency [12].

VII. CONCLUSIONS AND FURTHER DEVELOPMENT

An evaluation framework spanning over quantitative and qualitative techniques give a broader evaluation basis for music generated by machine learning based on the transformer architecture. Utilizing this framework to evaluate two different approaches for generating Americana music gave confidence in the superiority of a transfer learning based model (trained on the MAESTRO dataset) over a model solely trained on the Americana dataset.

Resulting songs generated by three-bar Americana primers was to some degree evaluated to be relatively close to human made music, with major differences between different melodies. However, the number of samples in the different evaluation approaches was relatively small. To increase the confidence in the results, a further work will be to massively scale up, both the number of generated songs for the the technical evaluation and the number of persons in the qualitative analysis.

We think that interviews with composers and musician about what makes "good" music is warranted. Moreover, we find it interesting to present AI-generated music to musical experts in the manner we have done to be particularly beneficial. An extension of this venue of research to consist of individual interviews and also contain a sort of "musical Turing test" is natural. As mentioned in chapter V, some form of systematic interviews has been utilized by Sturm & Ben-Tal [16], but when we combine this approach with the other 2 evaluation metrics described in this paper, the resulting framework becomes novel and more robust as an evaluation method for machine learning methods in the creative arts domain, such as music.

REFERENCES

1. Vaswani, A. *et al. Attention is all you need* in *Advances in neural information processing systems* (2017), 5998–6008.
2. Huang, C.-Z. A. *et al. Music transformer: Generating music with long-term structure* in *International Conference on Learning Representations* (2018).
3. Hawthorne, C. *et al. Enabling Factorized Piano Music Modeling and Generation with the MAESTRO Dataset* in *International Conference on Learning Representations* (2019). https://openreview.net/forum?id=r1lYRjC9F7.
4. Wu, S.-L. & Yang, Y.-H. The Jazz Transformer on the Front Line: Exploring the Shortcomings of AI-composed Music through Quantitative Measures. *arXiv preprint arXiv:2008.01307* (2020).
5. Eck, D. & Schmidhuber, J. *Finding temporal structure in music: Blues improvisation with LSTM recurrent networks* in *Proceedings of the 12th IEEE workshop on neural networks for signal processing* (2002), 747–756.
6. Huang, Y.-S. & Yang, Y.-H. Pop music transformer: Generating music with rhythm and harmony. *arXiv preprint arXiv:2002.00212* (2020).
7. Giraud, M., Groult, R., Leguy, E. & Levé, F. Computational fugue analysis. *Computer Music Journal* **39,** 77–96 (2015).
8. Hennion, A. The production of success: an antimusicology of the pop song. *Popular Music* **3,** 159–193 (1983).
9. Jazz., F. *Oxford Languages* Accessed 2 October 2020 (Oxford University Press, September 2020).

10. Center, A. F. *Folk Music and Song* Accessed 2 October 2020. %5Curl%7Bhttps://www.loc.gov/folklife/guide/folkmusicandsong.html%7D.

11. Dai, Z. *et al. Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context* 2019. arXiv: 1901.02860 `[cs.LG]`.

12. Kitaev, N., Kaiser, Ł. & Levskaya, A. *Reformer: The Efficient Transformer* 2020. arXiv: 2001.04451 `[cs.LG]`.

13. Donahue, C., Mao, H. H., Li, Y. E., Cottrell, G. W. & McAuley, J. LakhNES: Improving multi-instrumental music generation with cross-domain pre-training. *arXiv preprint arXiv:1907.04868* (2019).

14. Gan, C., Huang, D., Chen, P., Tenenbaum, J. B. & Torralba, A. *Foley Music: Learning to Generate Music from Videos* 2020. arXiv: 2007.10984 `[cs.CV]`.

15. Yang, L.-C. & Lerch, A. On the evaluation of generative models in music. *Neural Computing and Applications* **32,** 4773–4784 (2020).

16. Sturm, B. L. & Ben-Tal, O. Taking the models back to music practice: Evaluating generative transcription models built using deep learning. *Journal of Creative Music Systems* **2** (2017).

17. Shin, H. *et al.* Deep Convolutional Neural Networks for Computer-Aided Detection: CNN Architectures, Dataset Characteristics and Transfer Learning. *IEEE Transactions on Medical Imaging* **35,** 1285–1298 (2016).

18. Mislevy, R. J., Steinberg, L. S. & Almond, R. G. Focus article: On the structure of educational assessments. *Measurement: Interdisciplinary research and perspectives* **1,** 3–62 (2003).

19. Mislevy, R. J., Haertel, G., Riconscente, M., Rutstein, D. W. & Ziker, C. in *Assessing model-based reasoning using evidence-centered design* 19–24 (Springer, 2017).

20. Shaw, A. *MusicAutobot* https://github.com/bearpelican/musicautobot. (accessed: 13.11.2020).